

MACHINE COMPREHENSION USING MATCH-LSTM AND ANSWER POINTER

Nikitha Pasya (B00830710)

Introduction

- An end-to-end neural architecture for the question answering task.
- The architecture is based on match-LSTM and Pointer Net, a sequence-to-sequence model to constrain the output tokens to be from the input sequences.
- Given a piece of text, which is referred to as a passage, and a question related to the passage. The goal is to identify a subsequence from the passage as the answer to the question.

SQuAD

- SQuAD provides a challenging testbed for evaluating machine comprehension algorithms.
- In SQuAD the answers do not come from a small set of candidate answers and they have variable lengths.
- The questions and answers in SQuAD were created by humans through crowdsourcing, which makes the dataset more realistic.

In 1870, Tesla moved to Karlovac, to attend school at the Higher Real Gymnasium, where he was profoundly influenced by a math teacher **Martin Sekulic'**. The classes were held in **German**, as it was a school within the Austro-Hungarian Military Frontier. Tesla was able to perform integral calculus in his head, which prompted his teachers to believe that he was cheating. He finished a four-year term in three years, graduating in 1873.

- | | |
|--|---|
| 1. In what language were the classes given? | German |
| 2. Who was Tesla's main influence in Karlovac? | Martin Sekulic' |
| 3. Why did Tesla go to Karlovac? | attend school at the Higher
Real Gymnasium |

MATCH-LSTM

- To predict whether the premise(question) entails the hypothesis(passage), the match-LSTM model goes through the tokens of the hypothesis sequentially.
- At each position of the hypothesis, attention mechanism is used to obtain a weighted vector representation of the premise.
- Weighted premise is then to be combined with a vector representation of the current token of the hypothesis and fed into an LSTM, which we call the match-LSTM.
- Match-LSTM sequentially aggregates the matching of the attention-weighted premise to each token of the hypothesis and uses the aggregated matching result to make a final prediction.

Ptr-Net

- Ptr-Net uses attention mechanism as a pointer to select a position from the input sequence as an output symbol.
- To construct answers using tokens from the input text.
- Two ways of using Pointer Net for QA.
 - Sequential model
 - Boundary model

Overview of Neural network model

- The passage is represented by matrix $P \in R^{d \times P}$ and the question is represented by matrix $Q \in R^{d \times Q}$

- **LSTM Preprocessing Layer**

standard one-directional LSTM to preprocess the passage and the question separately

$$H^p = \overrightarrow{LSTM}(P), H^q = \overrightarrow{LSTM}(Q)$$

- **Match-LSTM Layer**

- tries to match the passage against the question.
- At position i of the passage, it first uses the standard word-by-word attention mechanism to obtain attention weight vector $\overrightarrow{\alpha}_i \in R^Q$ as follows:

$$\overrightarrow{G}_i = \tanh(W^q H^q + (W^p h_i^p + W^r \overrightarrow{h}_{i-1}^r + b^p) \otimes e_Q)$$

$$\overrightarrow{\alpha}_i = \text{softmax}(w^T \overrightarrow{G}_i + b \otimes e_Q)$$

use the attention weight vector $\overrightarrow{\alpha_i}$ to obtain a weighted version of the question and combine it with the current token of the passage to form a vector $\overrightarrow{z_i}$.

- $\overrightarrow{z_i}$ is fed into a standard one-directional LSTM to form match-LSTM:

$$\overrightarrow{h_i^r} = \overrightarrow{LSTM}(\overrightarrow{z_i}, \overrightarrow{h_{i-1}^r})$$

$$\text{where } \overrightarrow{h_i^r} \in R^{(l)}$$

- further build a similar match-LSTM in the reverse direction

$$\overleftarrow{G_i} = \tanh(W^q H^q + (W^p h_i^p + W^r \overleftarrow{h_{i+1}^r} + b^p) \otimes e_q)$$

$$\overleftarrow{\alpha_i} = \text{softmax}(w^T \overleftarrow{G_i} + b \otimes e_q)$$

- Let $\overrightarrow{H^r} \in R^{(l * P)}$ represent the hidden states $[\overrightarrow{h_1^r}, \overrightarrow{h_2^r}, \overrightarrow{h_3^r}, \dots, \overrightarrow{h_p^r}]$ and $\overleftarrow{H^r} \in R^{(l * P)}$ represents $[\overleftarrow{h_1^r}, \overleftarrow{h_2^r}, \overleftarrow{h_3^r}, \dots, \overleftarrow{h_p^r}]$

$$H^r \in R^{(2l * P)} \text{ is the concatenation of two } \begin{bmatrix} \overrightarrow{H^r} \\ \overleftarrow{H^r} \end{bmatrix}$$

- **Answer Pointer Layer**

- uses Ptr-Net to select a set of tokens from the passage as the answer.
- This layer uses the sequence H^r as input.

The Sequence Model: The answer is represented by a sequence of integers $a = (a_1, a_2, \dots)$.

- Allows each a_k to take up an integer value between 1 and $P + 1$, Once $a_k = P + 1$, the generation of the answer stops.
- To generate the answer token a_k , the attention mechanism is used again to obtain an attention weight vector $\beta_k \in R^{(P+1)}$, where $\beta_{k,j} (1 \leq j \leq P + 1)$ is the probability of selecting the j th token from the passage as the k th token in the answer, β_k is modeled as follows:

$$F_k = \tanh(V \tilde{H}^r + (W^a h_{k-1}^a + b^a) \otimes e_{(P+1)}),$$

$$\beta_k = \text{softmax}(V^T F_k + c \otimes e_{(P+1)})$$

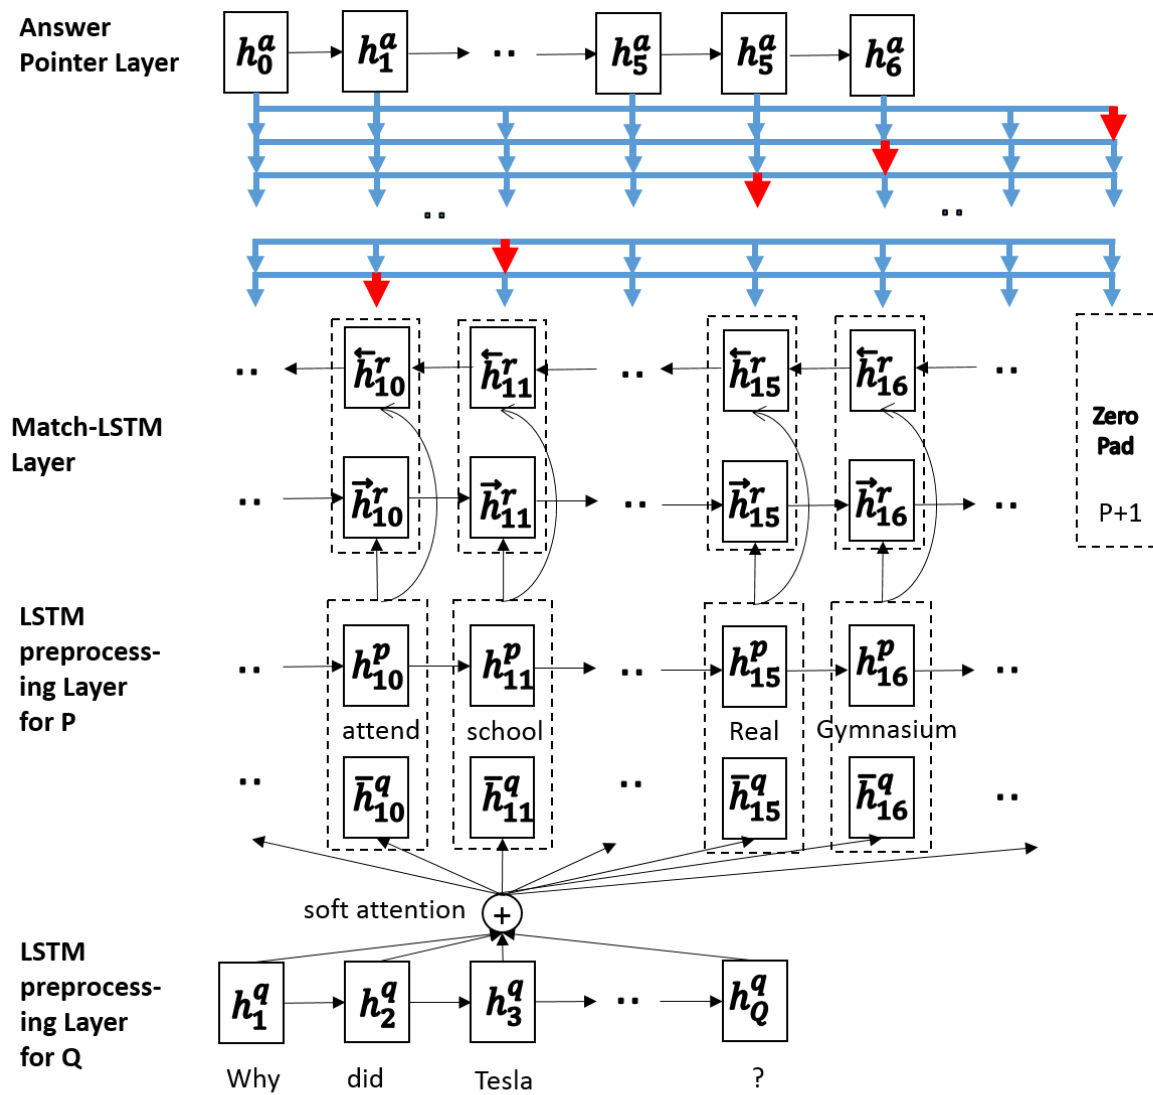
Where $\tilde{H}^r = [H^r; 0]$

- We can then model the probability of generating the answer sequence as
 $p(a | H^r) = \prod_k p(a_k | a_1, a_2, \dots, a_{k-1}, H^r)$
 $p(a_k = j | a_1, a_2, \dots, a_{k-1}, H^r) = \beta_{k,j}$

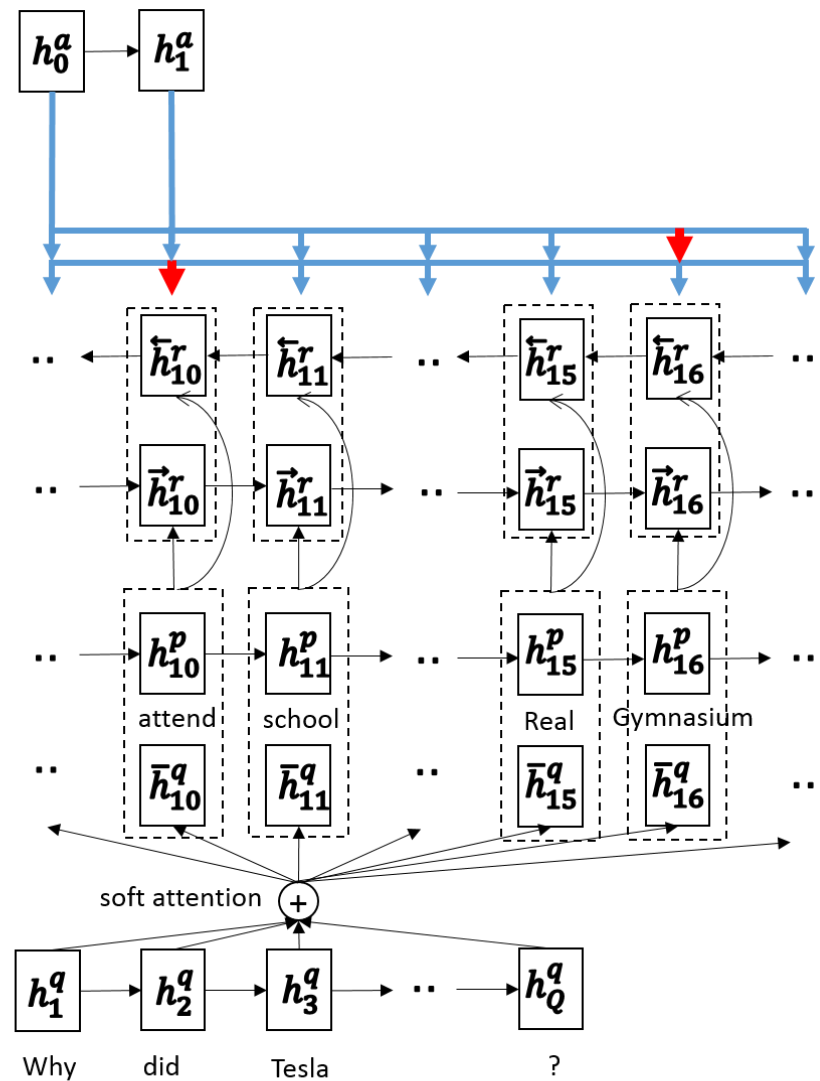
- **The Boundary Model:** Predicts two indices a_s and a_e .
 - the probability of generating an answer is simply modeled as

$$p(a | H^r) = p(a_s | H^r) p(a_e | a_s, H^r)$$

- the boundary model is further extended by incorporating a search mechanism. During prediction, the length of the span is limited and globally search the span with the highest probability computed by $p(a_s) \times p(a_e)$.
- As the boundary has a sequence of fixed number of values, bi-directional Ans-Ptr can be simply combined to fine-tune the correct span.



(a) Sequence Model



(b) Boundary Model

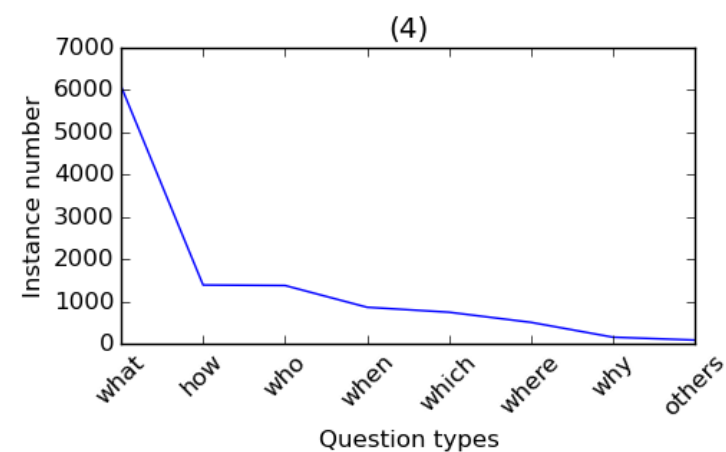
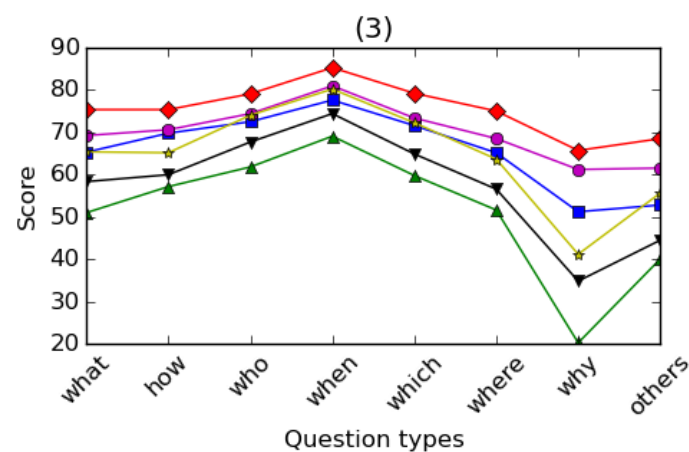

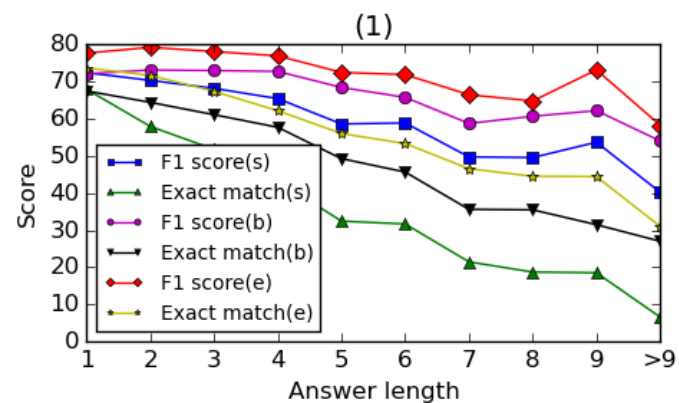
Experiments

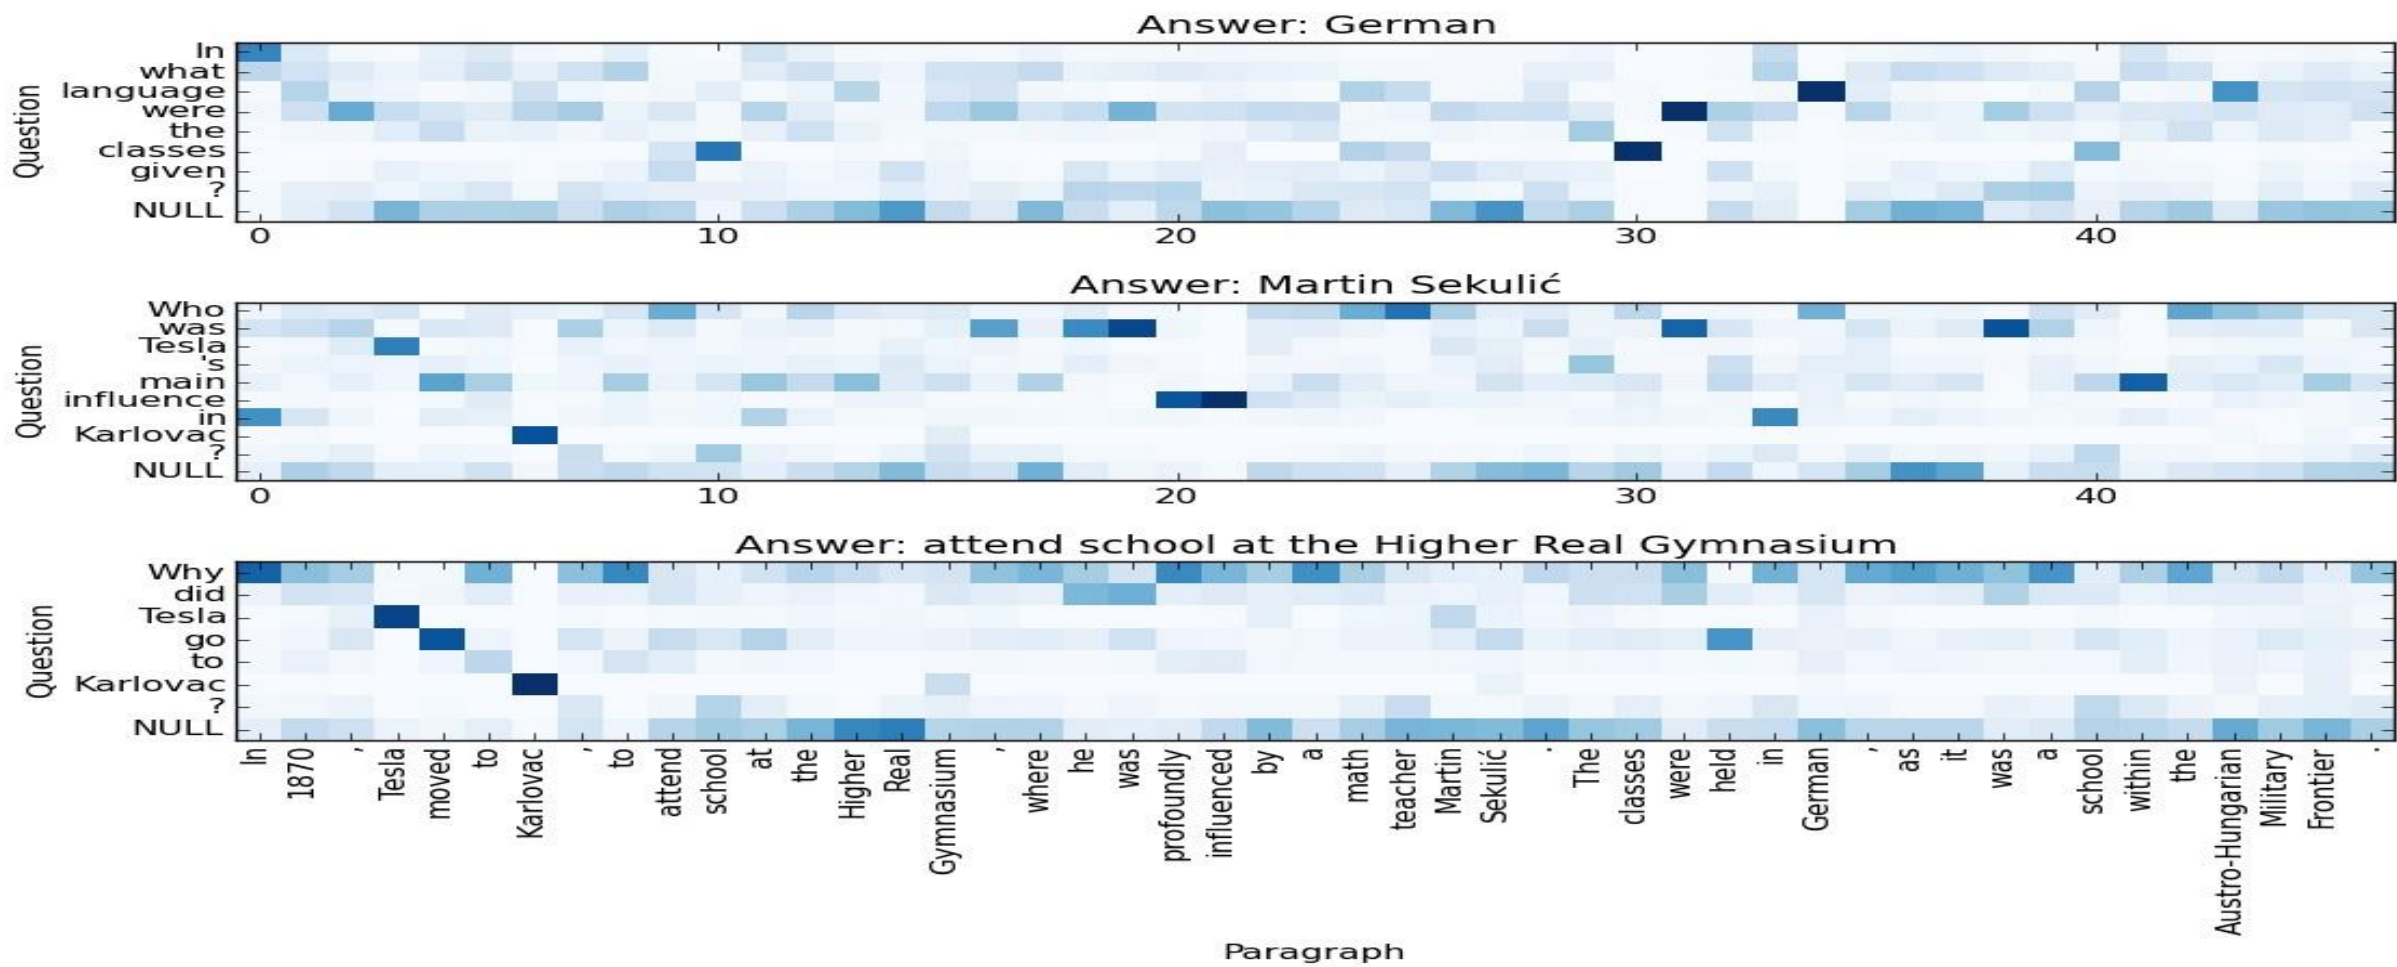
- The dimensionality l of the hidden layers is set to be 150 or 300.
- We use ADAMAX with the coefficients $\beta_1 = 0.9$ and $\beta_2 = 0.999$ to optimize the model. Each update is computed through a minibatch of 30 instances.
- The performance is measured by two metrics: percentage of exact match with the ground truth answers, and word-level F1 score when comparing the tokens in the predicted answers with the tokens in the ground truth answers.
- F1 scores with the best matching answers are used to compute the average F1 score.

Results

	l	$ \theta $	Exact Match		F1	
			Dev	Test	Dev	Test
Random Guess	-	0	1.1	1.3	4.1	4.3
Logistic Regression	-	-	40.0	40.4	51.0	51.0
DCR	-	-	62.5	62.5	71.2	71.0
Match-LSTM with Ans-Ptr (Sequence)	150	882K	54.4	-	68.2	-
Match-LSTM with Ans-Ptr (Boundary)	150	882K	61.1	-	71.2	-
Match-LSTM with Ans-Ptr (Boundary+Search)	150	882K	63.0	-	72.7	-
Match-LSTM with Ans-Ptr (Boundary+Search)	300	3.2M	63.1	-	72.7	-
Match-LSTM with Ans-Ptr (Boundary+Search+b)	150	1.1M	63.4	-	73.0	-
Match-LSTM with Bi-Ans-Ptr (Boundary+Search+b)	150	1.4M	64.1	64.7	73.9	73.7
Match-LSTM with Ans-Ptr (Boundary+Search+en)	150	882K	67.6	67.9	76.8	77.0

Further Analysis





Conclusion

- Developed two models for the machine comprehension problem defined in the Stanford Question Answering (SQuAD) dataset, both making use of match-LSTM and Pointer Network.
- Experiments on the SQuAD dataset showed that the boundary model, could achieve an exact match score of 67.6% and an F1 score of 77% on the test dataset.

References

- S.Wang and J.Jiang.Machine Comprehension using Match-LSTM and Answer pointer. arXiv preprint arXiv:1608.07905, 2017.
- Oriol vinyals, Meire Fortunato ,Navdeep Jaitly.Pointer Networks.arXiv preprint arxiv.org:1506.03134,2017.